


# SIMR: Single Instruction Multiple Request Processing for Energy-Efficient Data Center Microservices


**Mahmoud Khairy\***, Ahmad Alawneh, Aaron Barnes, and Timothy G. Rogers

Purdue University

# Datacenter Power Breakdown



**Datacenter Power Breakdown  
(from Google)**



**CPU Power Breakdown**

25-45% of datacenter power is consumed in CPU's instruction supply (frontend & OoO)

# 1 Application, Millions of Users

Google

facebook

Private Datacenter

Uber

NETFLIX

AWS

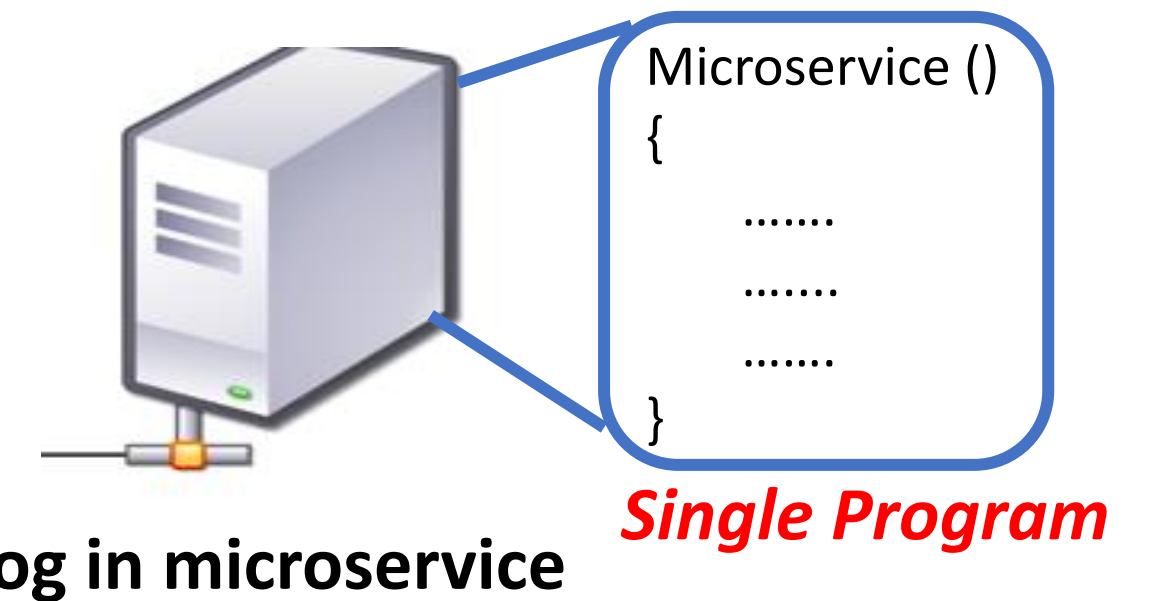
Public Datacenter

*“Similar” Request-Level Parallelism*

*1000s of independent requests are all running the same code*

Log-in reqs  
("xyz", "1234")  
("john", "5678")  
("ma98", "4444")  
("mah", "ko56")

*Multiple Data*



search reqs  
("purdue univ")  
("arsenal fc")  
("elections 2024")  
("stock today")



**Key Observation #1:** Single Program Multiple Data (SPMD) are abundant in the datacenters

# Server Workloads on GPU's

- **Key Idea:** Exploit SPMD by batching requests and run them on GPU's Single Instruction Multiple Thread (SIMT) or CPU's SIMD
- **Advantage:** Significant energy efficiency (throughput/watts) vs multi-threaded CPU
- **Drawbacks:**
  - (1) Hindering programmability (C++/PHP vs CUDA/OpenCL)
  - (2) Limited system calls support
  - (3) High service latency (10-6000x)
    - GPUs tradeoff single threaded optimizations (OoO, speculative execution, etc.) in favor of excessive multithreading
    - In SIMD, relying on branch predicates & fine grain context

## Rhythm: Harnessing Data Parallel Hardware for Server Workloads

Sandeep R Agrawal  
Duke University  
sandeep@cs.duke.edu

Valentin Pistol  
Duke University  
pistol@cs.duke.edu

Jun Pang  
Duke University  
pangjun@cs.duke.edu

John Tran  
NVIDIA  
johntran@nvidia.com

David Tarjan\*  
NVIDIA

Alvin R Lebeck  
Duke University  
alvy@cs.duke.edu

Rhythm, ASPLOS 2014

## MemcachedGPU: Scaling-up Scale-out Key-value Stores

Tayler H. Hetherington  
The University of British Columbia  
taylerh@ece.ubc.ca

Mike O'Connor  
NVIDIA & UT-Austin  
moconnor@nvidia.com

Tor M. Aamodt  
The University of British Columbia  
aamodt@ece.ubc.ca

MemcachedGPU, SoCC 2015

## ispc: A SPMD Compiler for High-Performance CPU Programming

Matt Pharr  
Intel Corporation  
matt.pharr@intel.com

William R. Mark  
Intel Corporation  
william.r.mark@intel.com

ispc, InPar 2012

Recall: GPUs and SIMDs were designed to execute data parallel portion (i.e., loops) not the entire application

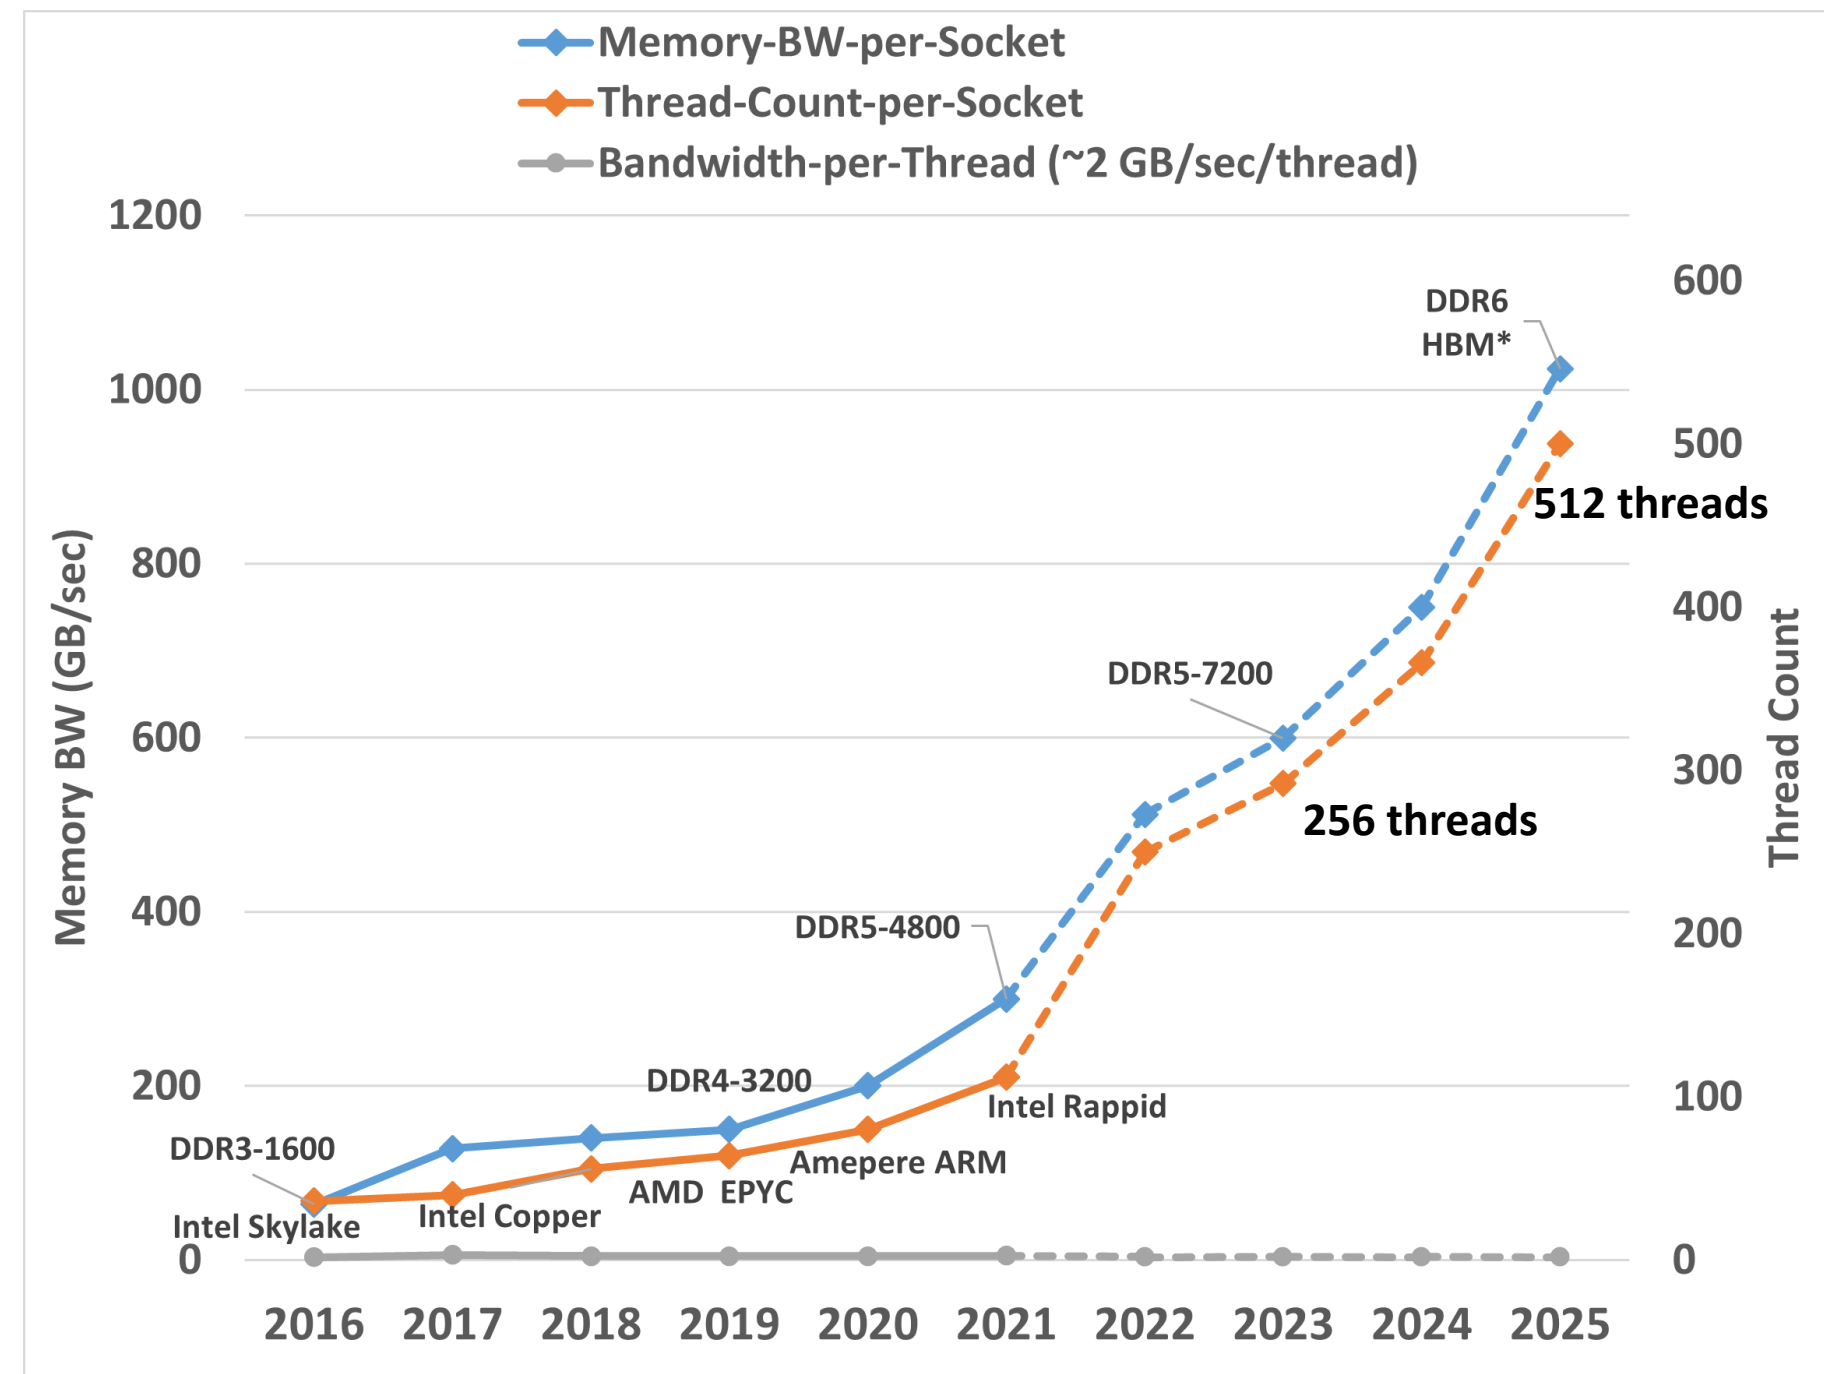
*“Slower but energy-efficient wimpy cores only win for general data center workloads if their single-core speed is reasonably close to that of mid-range brawny cores”*

*Up to 2x slower latency can be tolerated by data center providers*



**Urs Hölzle**  
**Google SVP**



# Off-Chip BW Scaling



**Key Observation #2:** There is available headroom to increase on-chip throughput (thread count) in the foreseeable future.



# How to increase on-chip throughput of CPU?

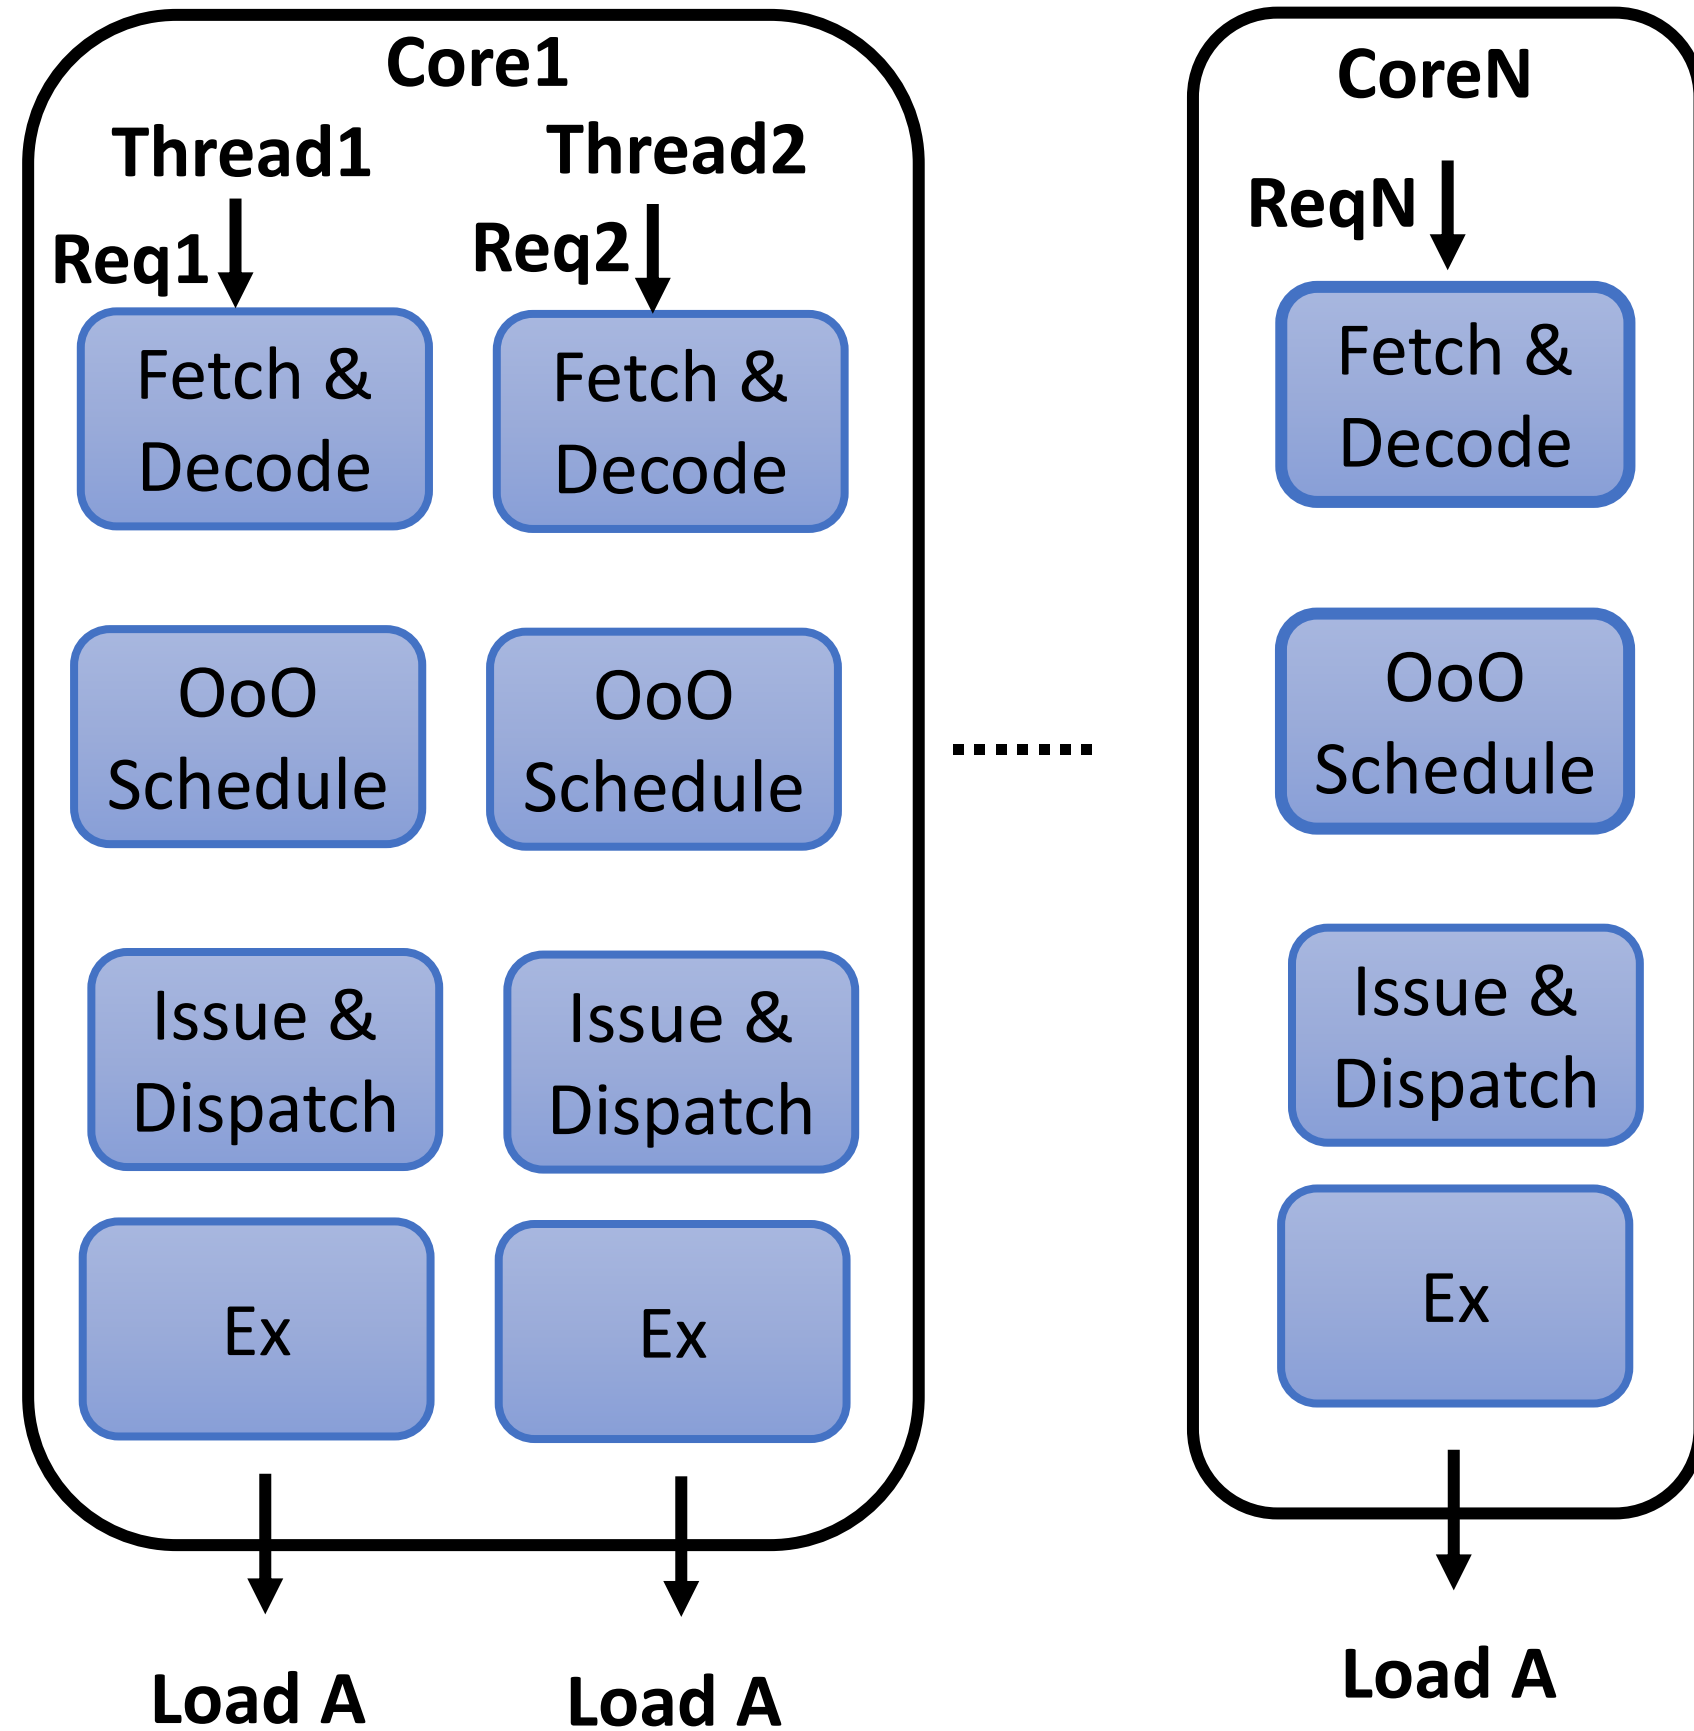
- Direction#1 (industry standard): Add more Chiplets + Cores + SMT 
- Direction#2 (this work): Move to *SIMT* 
  - More energy efficient (throughput/watts)
  - Cost-effective (throughput/area)
  - Better scalability

*“Let’s bring SIMD efficiency to the CPU world!”*




# SIMT Efficiency

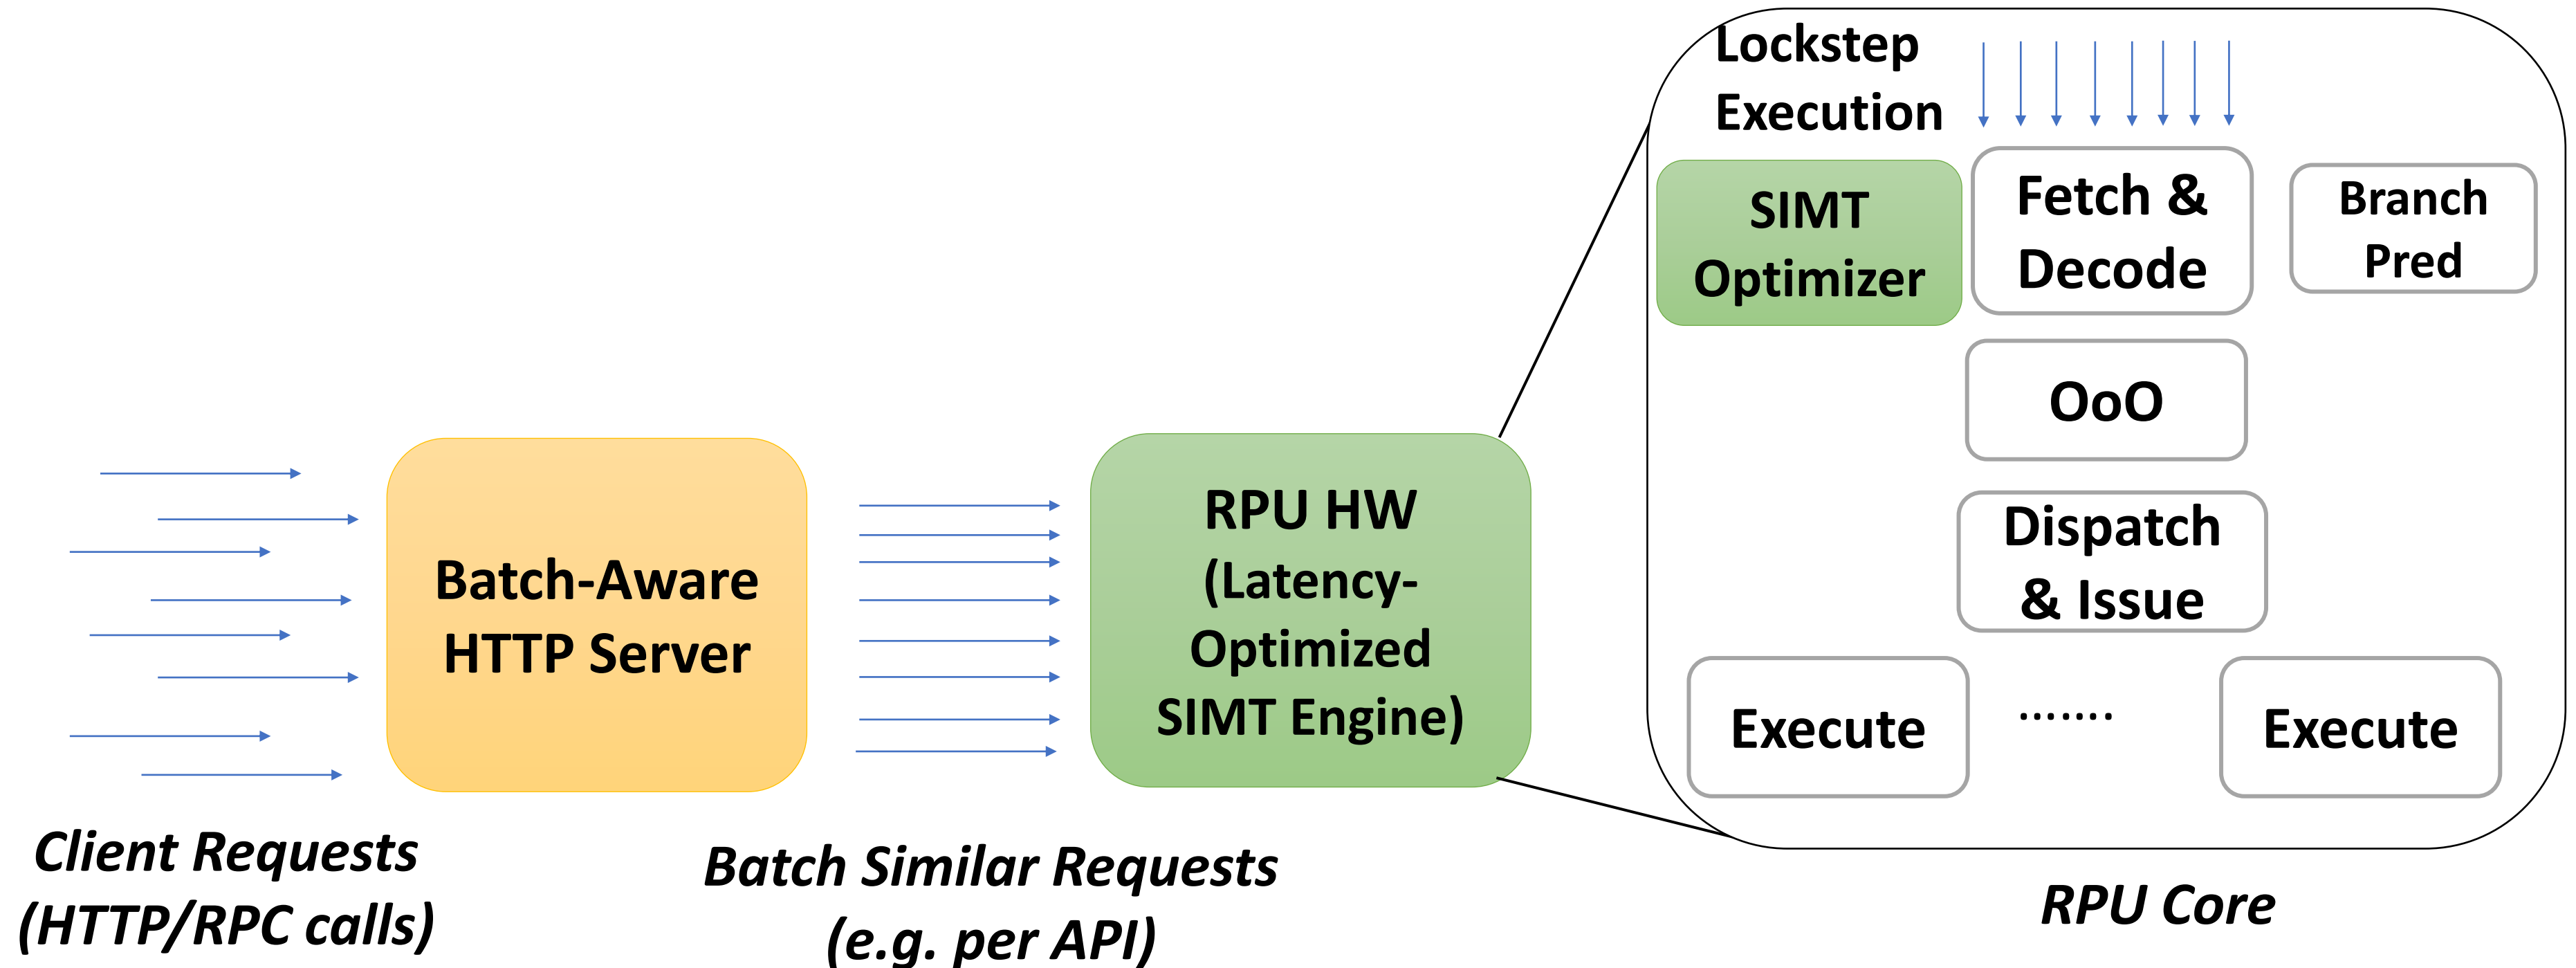
CPU Multi-Core with Simultaneous Multi-Threading



Request Processing Unit (RPU)  
SIMT+OoO Architecture



# SIMR System Overview



# CPU vs GPU vs RPU

Metric	CPU	GPU	RPU
Core model	OoO	In-Order	OoO
Programming	General-Purpose	CUDA/OpenCL	General-Purpose
ISA	x86/ARM	HSAIL/PTX	x86/ARM
System Calls Support	Yes	No	Yes
Thread grain	Coarse grain	Fine grain	Coarse grain
Threads per core	Low (1-8)	Massive (2K)	Moderate (8-32)
Thread model	SMT	SIMT	SIMT
Consistency	Variant	Weak+NMCA*	Weak+NMCA*
Interconnect	Mesh/Ring	Crossbar	Crossbar

The RPU takes advantage of the latency optimizations and programmability of the CPU

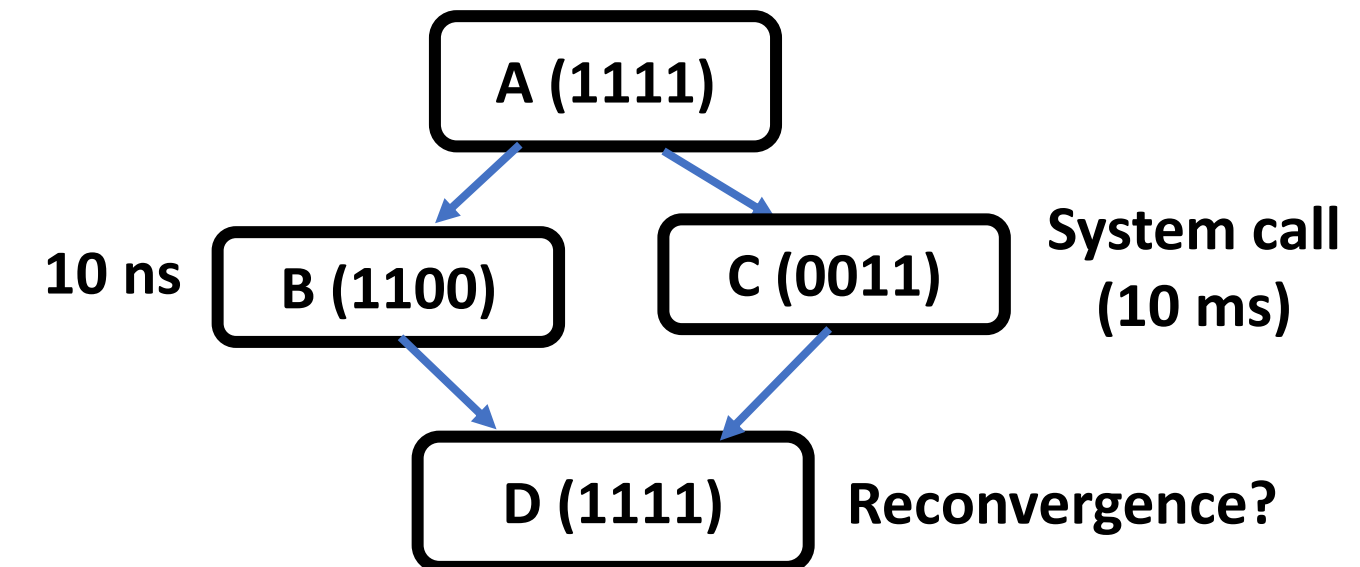
& SIMT efficiency and memory model scalability of the GPU

\*NMCA: non-multi copy atomicity

# RPU's Challenges

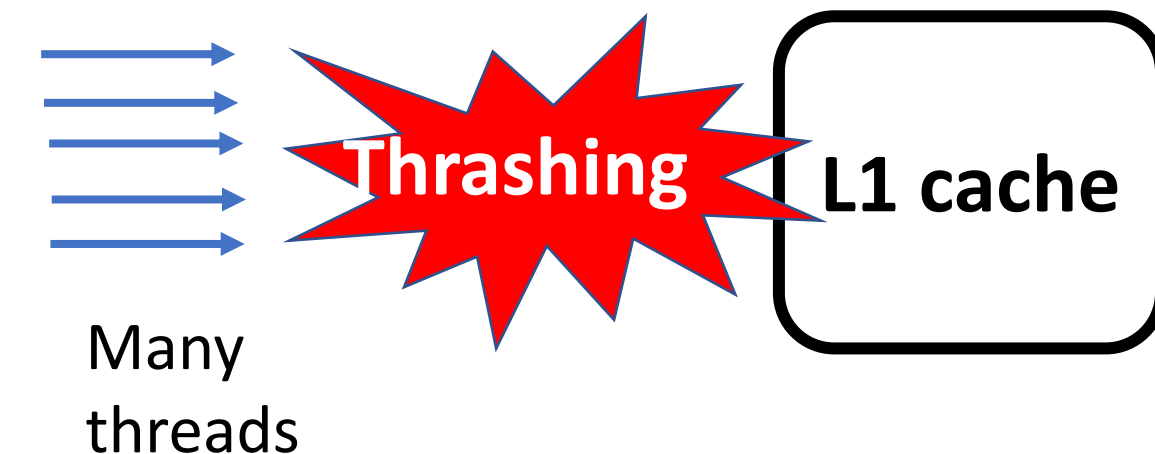
- Control Divergence

- Challenge: Control divergence with high latency path
- Solution: Optimized batching & System-level batch split



- Memory Divergence

- Challenge: Cache/TLB contention & bank conflicts
- Solution: Batch tuning, stack/memory coalescing and SIMR-aware memory allocation



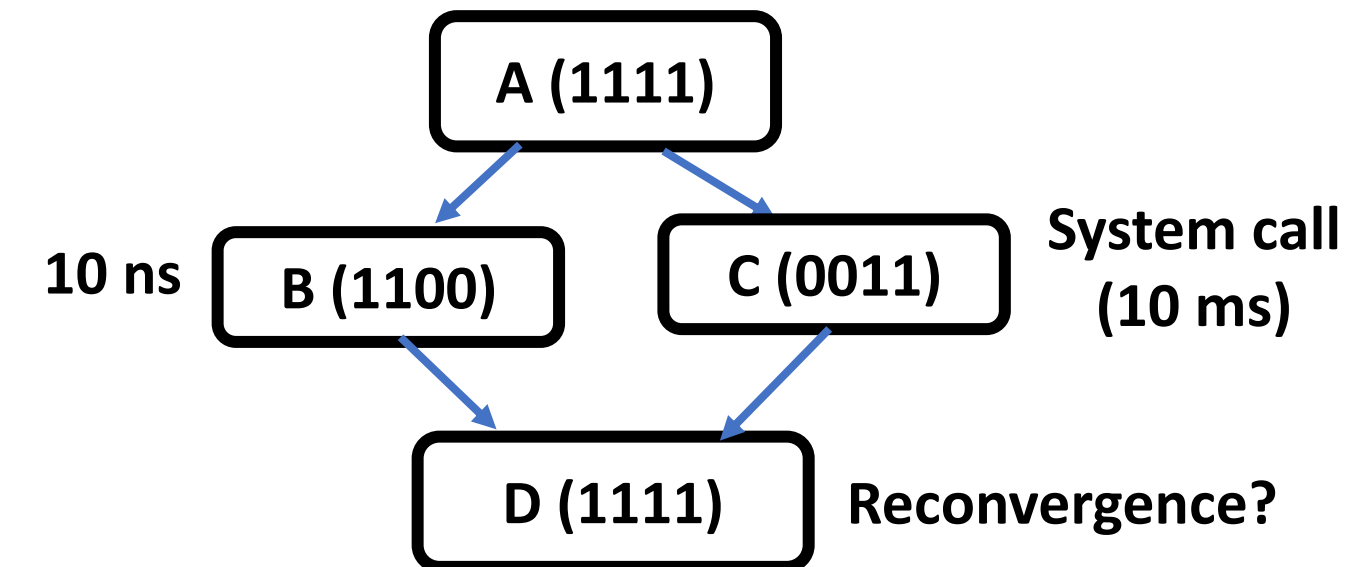
- Larger execution units & cache resources

- Challenge: Higher instruction execution & L1 hit latency
- Solution: Exploit low IPC, less generated traffic and employ sub-batching interleaving

# RPU's Challenges

- Control Divergence

- Challenge: Control divergence with high latency path
- Solution: Optimized batching & System-level batch split



- Memory

- Challenge: High memory access latency
- Solution: Batch tuning, stack/memory coalescing and SIMR-aware memory allocation

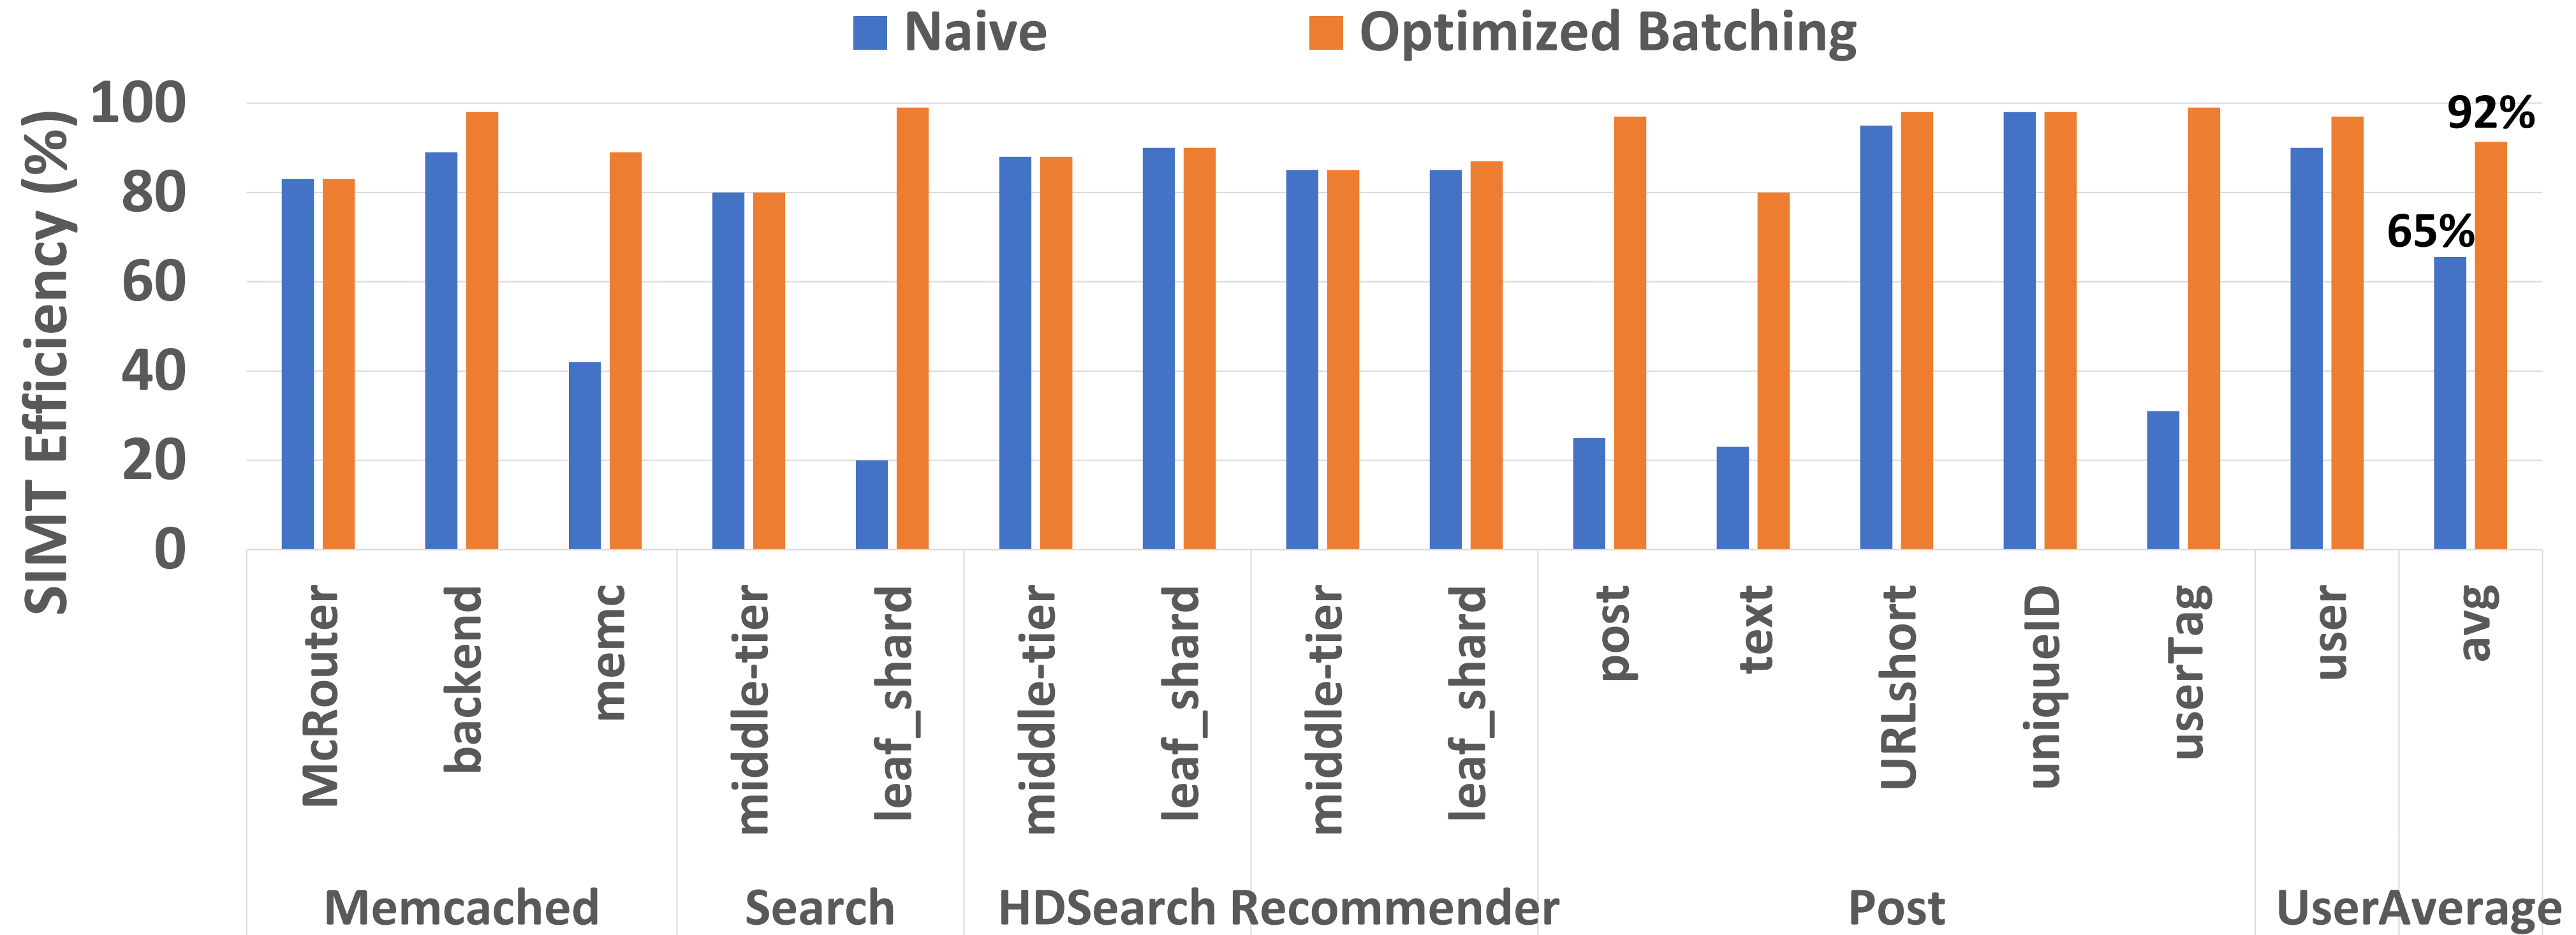
**Read more details in the paper on how we address these challenges**



- Larger execution units & cache resources


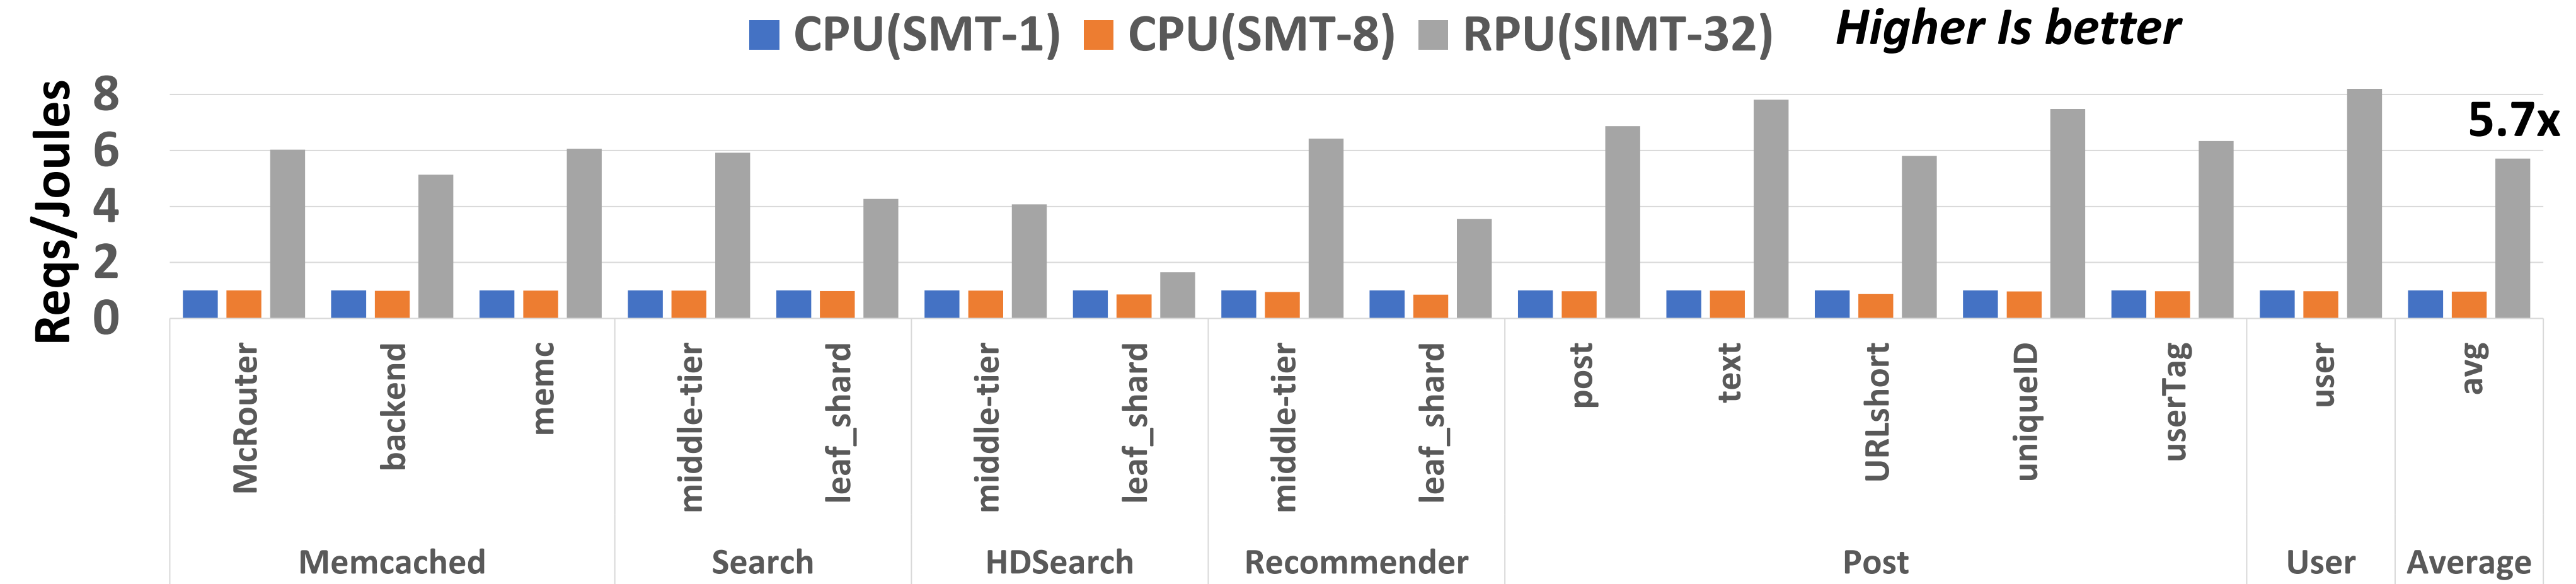
- Challenge: Higher instruction execution & L1 hit latency
- Solution: Exploit low IPC, less generated traffic and employ sub-batching interleaving

# SIMT Control Efficiency



Notes: (1) Batch Size = 32 & #batches=75, (2) System Calls are not traced, (3)  $\text{SIMT Eff} = \text{scalar-instructions} / (\text{batch-instructions} * \text{batch-size})$ , (4) fine-grain locking are assumed. Other assumptions are included in the paper.


# Efficiency and Service Latency Results (Simulation)






# Efficiency and Service Latency Results (Simulation)

CPU(SMT-1) CPU(SMT-8) RPU(SIMT-32) *Higher Is better*



RPU system achieves 5.7x higher throughput/Watt while maintaining acceptable service latency, without changing the programming SW interface



# Summary

- Request Similarity is abundant in the data center.
- We start with OoO CPU design and augment it with SIMT execution to maximize chip utilization and exploit the similarity.
- We co-design the software stack to support batching and awareness of SIMT execution.

*SIMT efficiency is high in the open-source microservices we study.*



**DeathStarBench**

$\mu$ Suite: A Benchmark Suite for Microservices

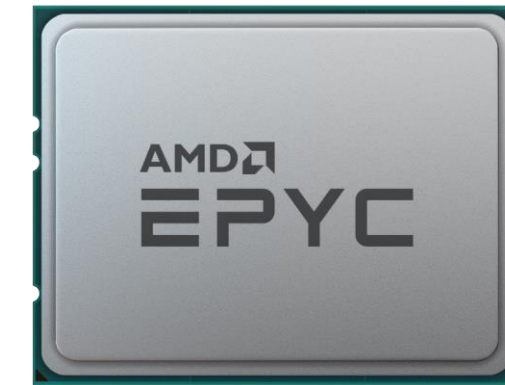
*We are very interested in evaluating SIMT control efficiency in proprietary production microservices.*

**Google**  
**facebook**

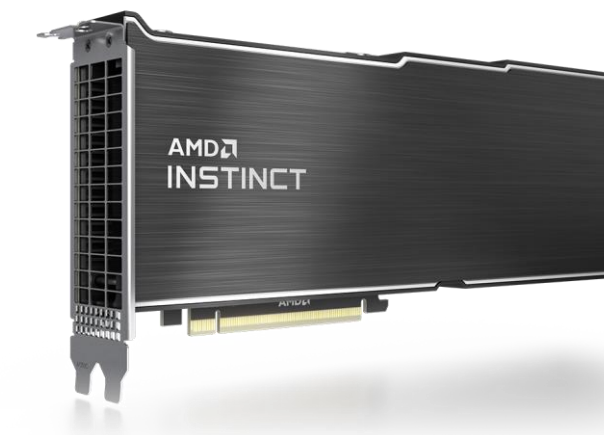
# *Thank You!*

## Q&A?


Instruction level parallelism (ILP) &  
Thread level parallelism (TLP)



Data level parallelism (DLP)

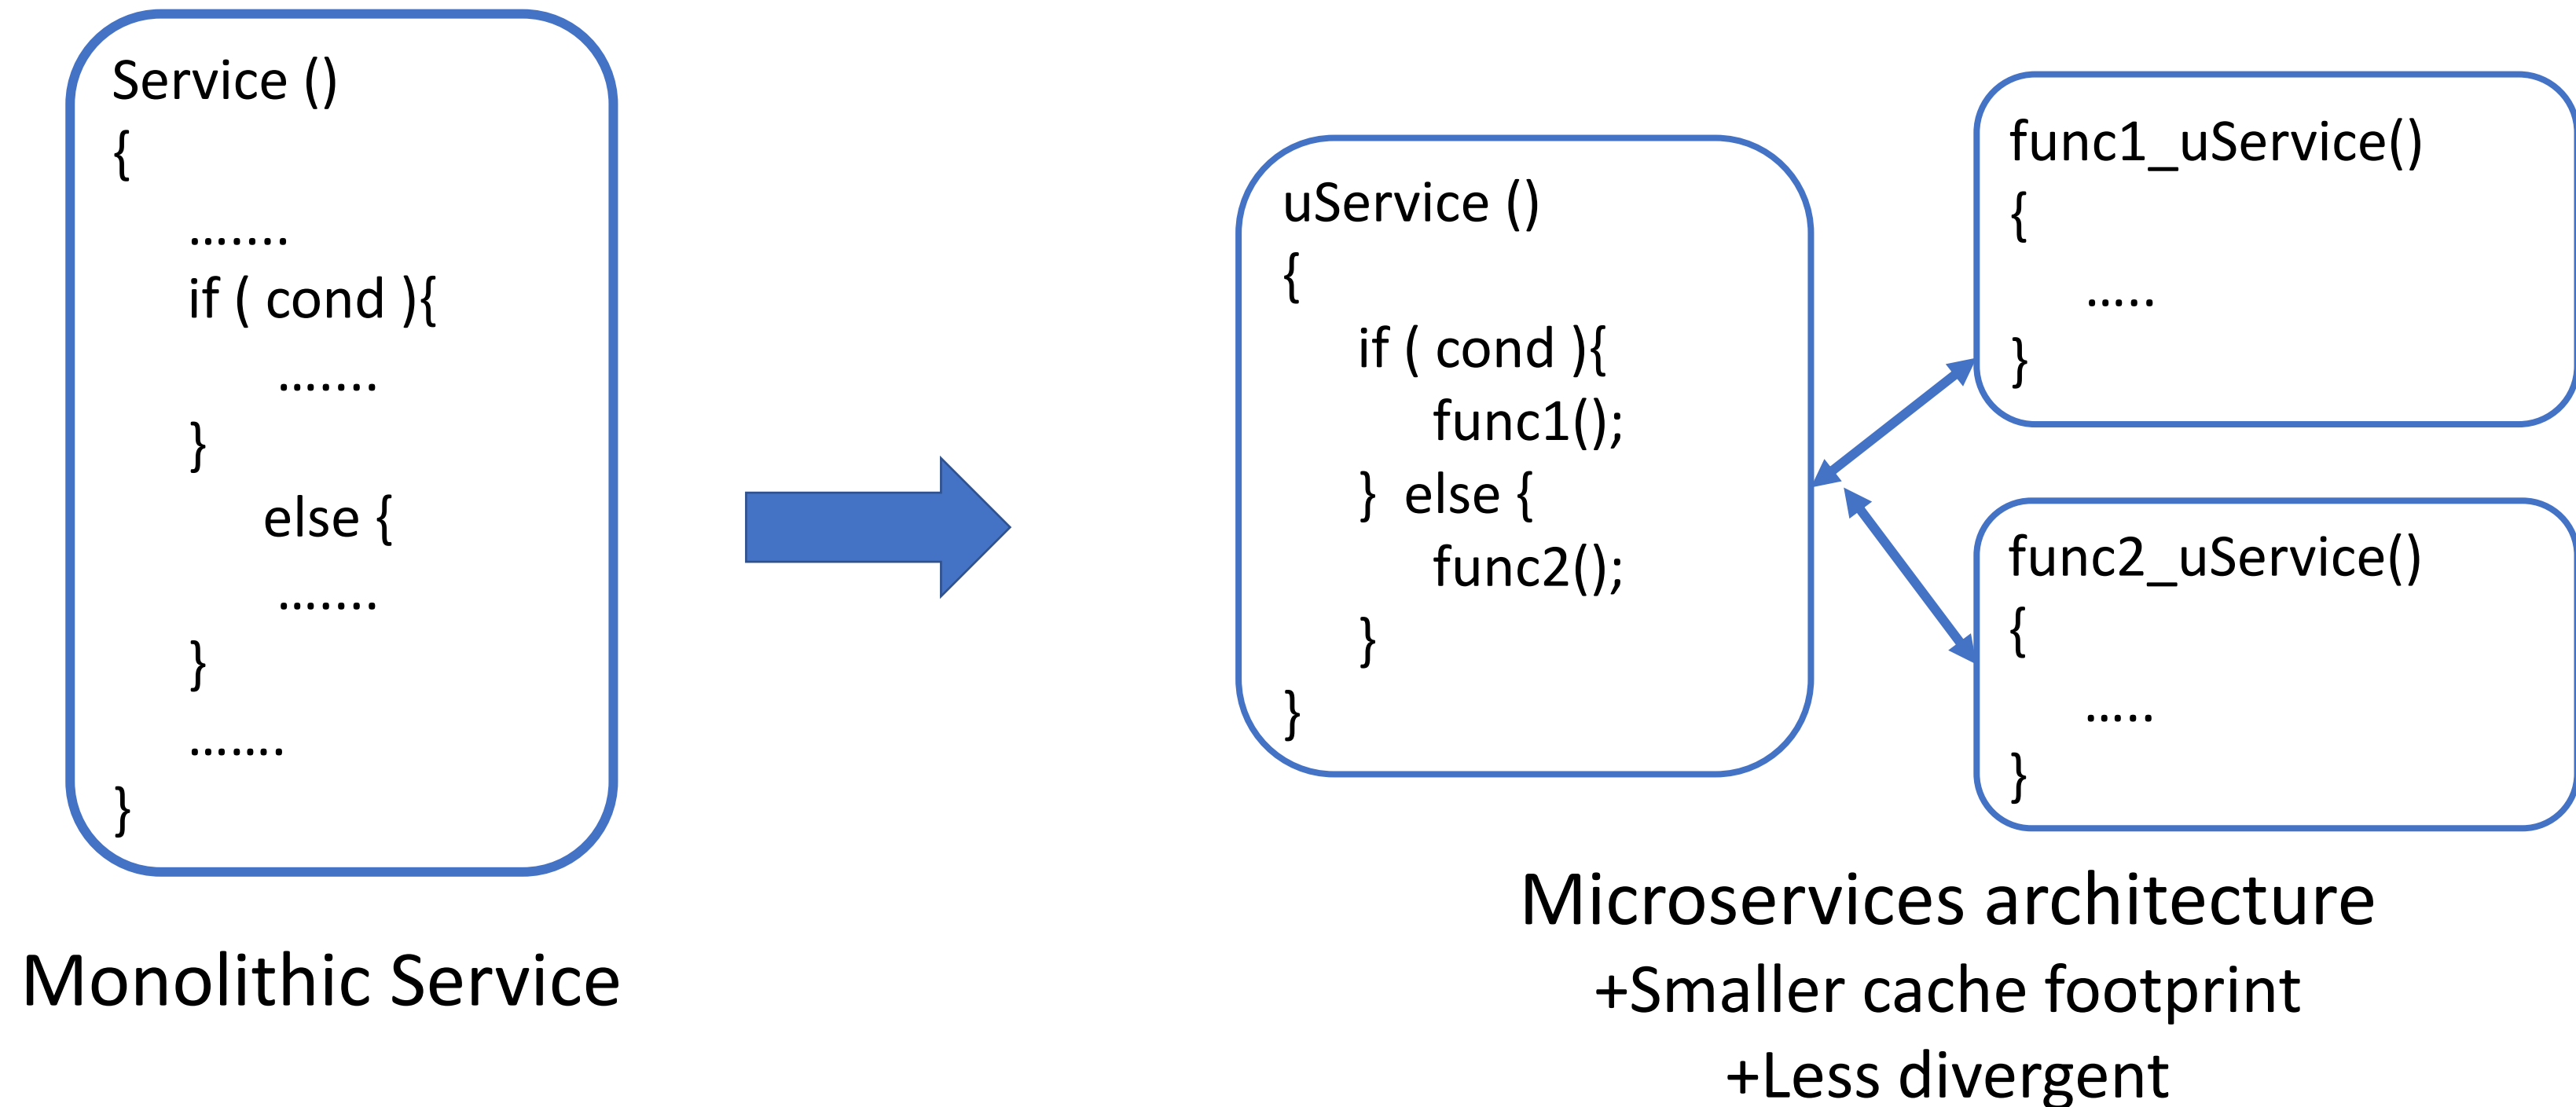


Request level parallelism (RLP)



# Backup Slides

# SIMT-friendly Microservices



**Key Observation#3:** Microservices reduce the per-thread cache requirement and minimize control-flow variations between concurrent threads

# Batching Optimization

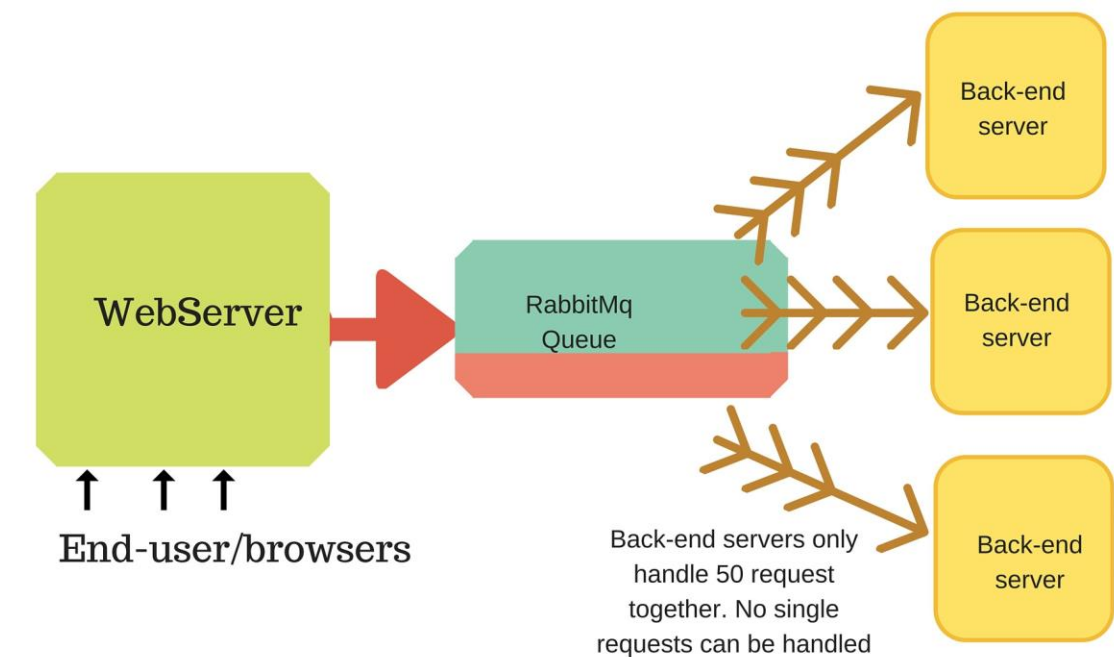
From Google's Production DL Inference

Production						MLPerf 0.7		
DNN	ms	batch	DNN	ms	batch	DNN	ms	batch
MLP0	7	200	RNN0	60	8	Resnet50	15	16
MLP1	20	168	RNN1	10	32	SSD	100	4
CNN0	10	8	BERT0	5	128	GNMT	250	16
CNN1	32	32	BERT1	10	64			

Table 5. Latency limit in ms and batch size picked for TPUv4i.


DL Inference Batching

Memcached servers



Network Batching

Power management

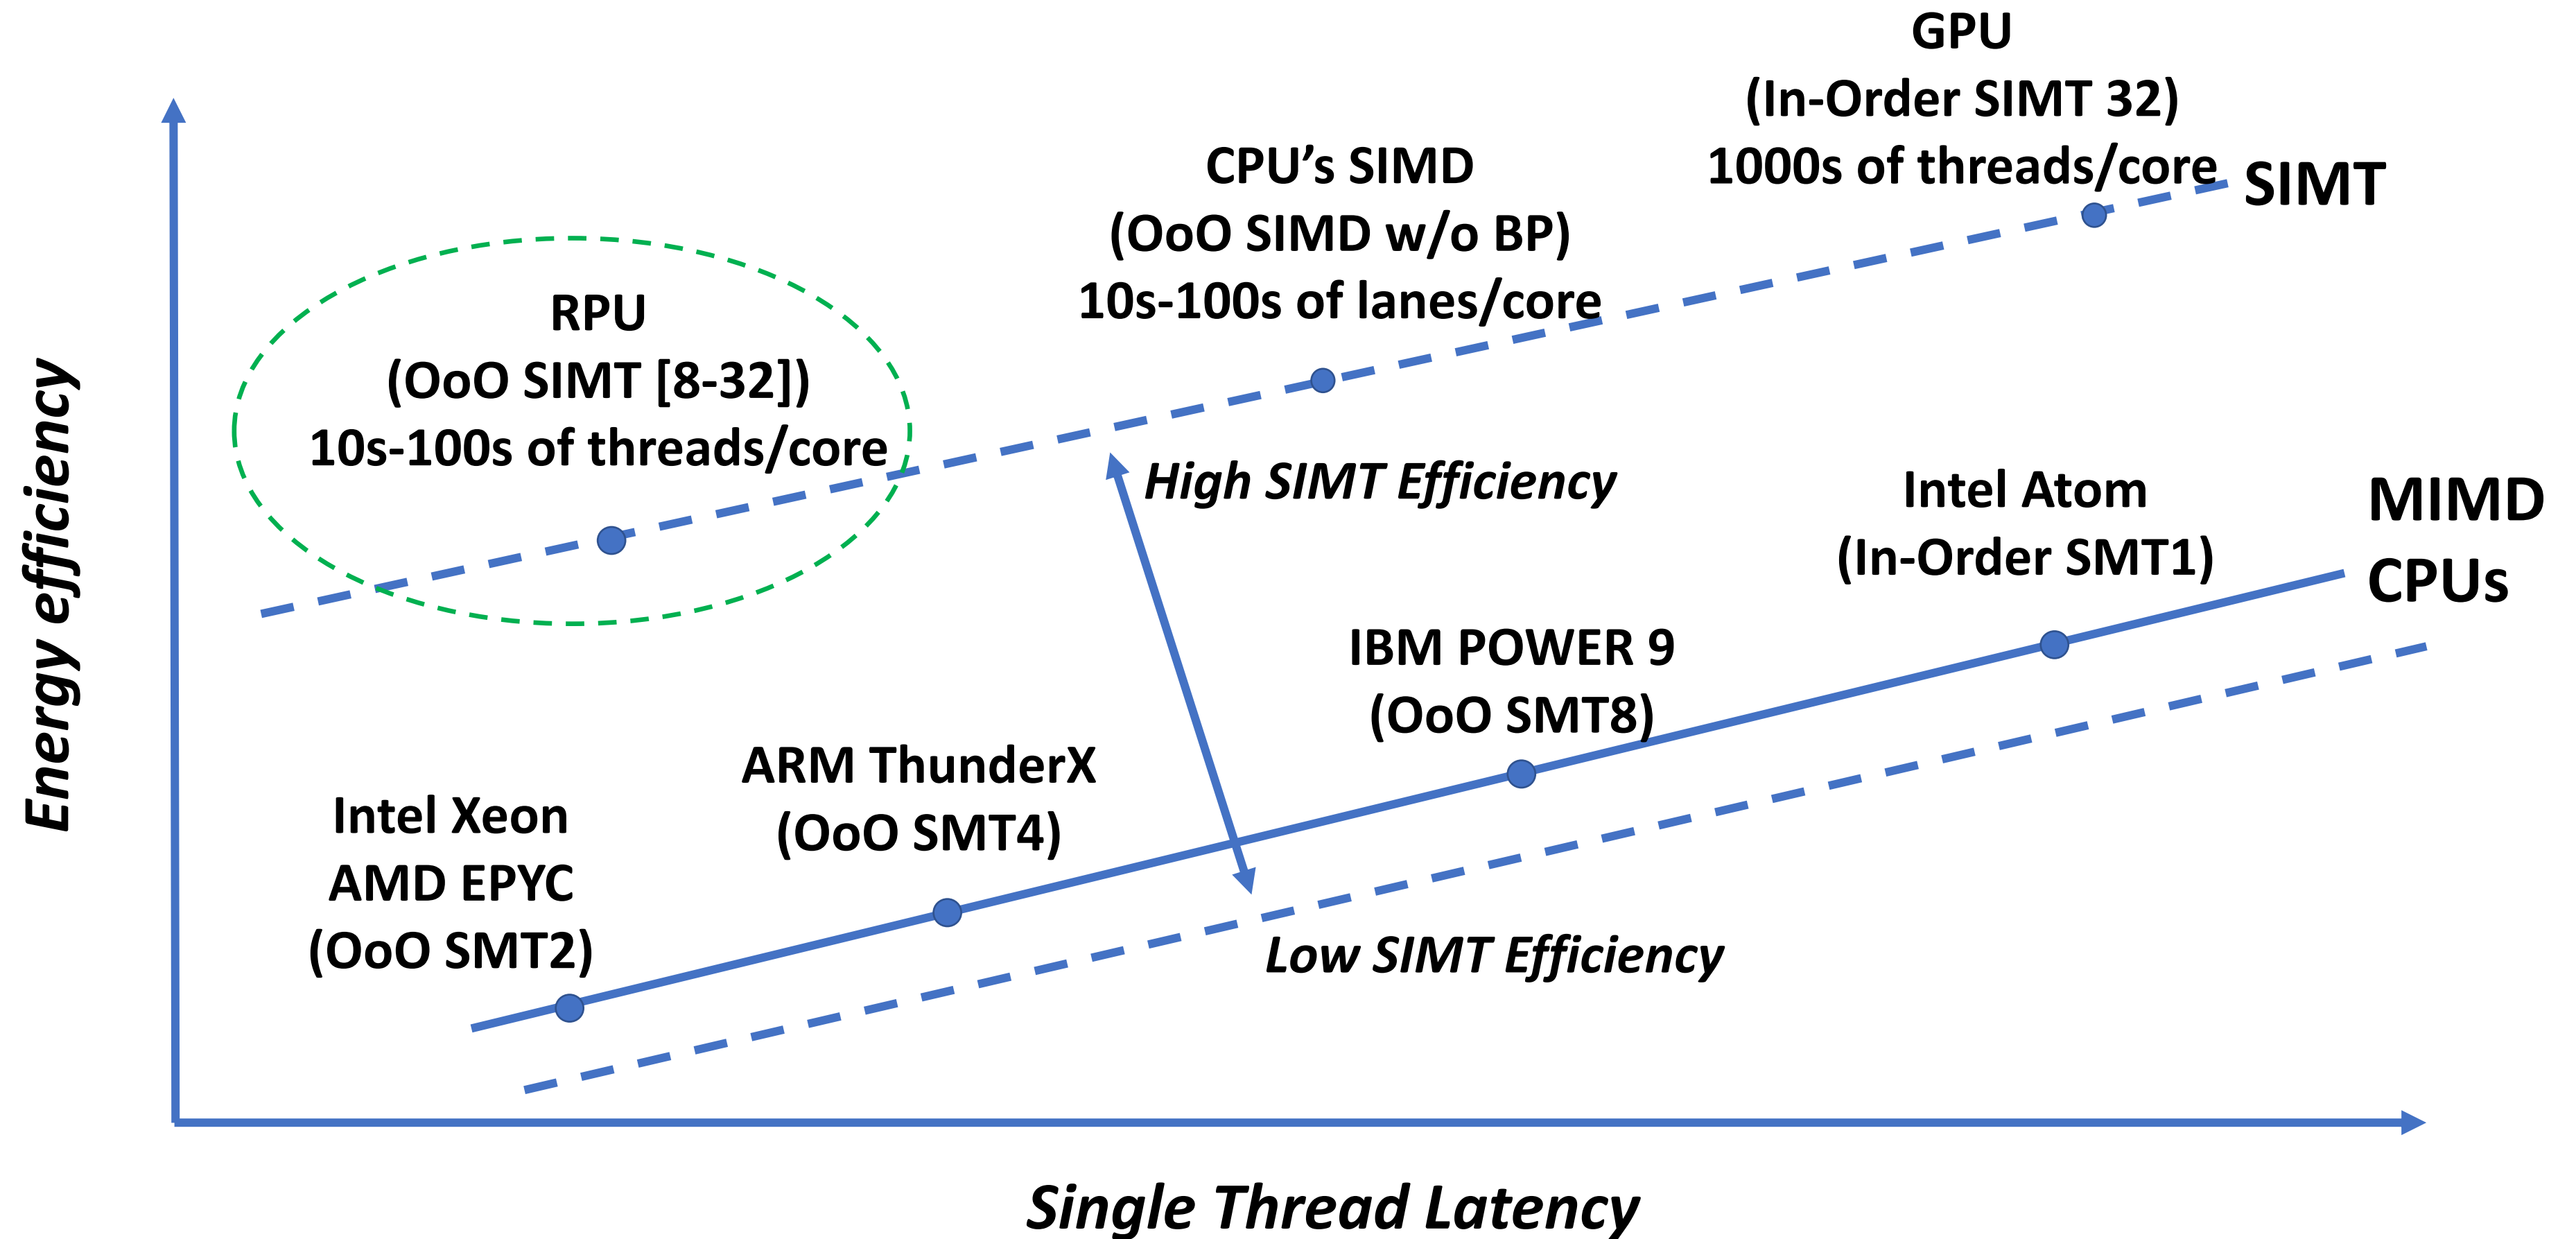


Batching for deep sleep

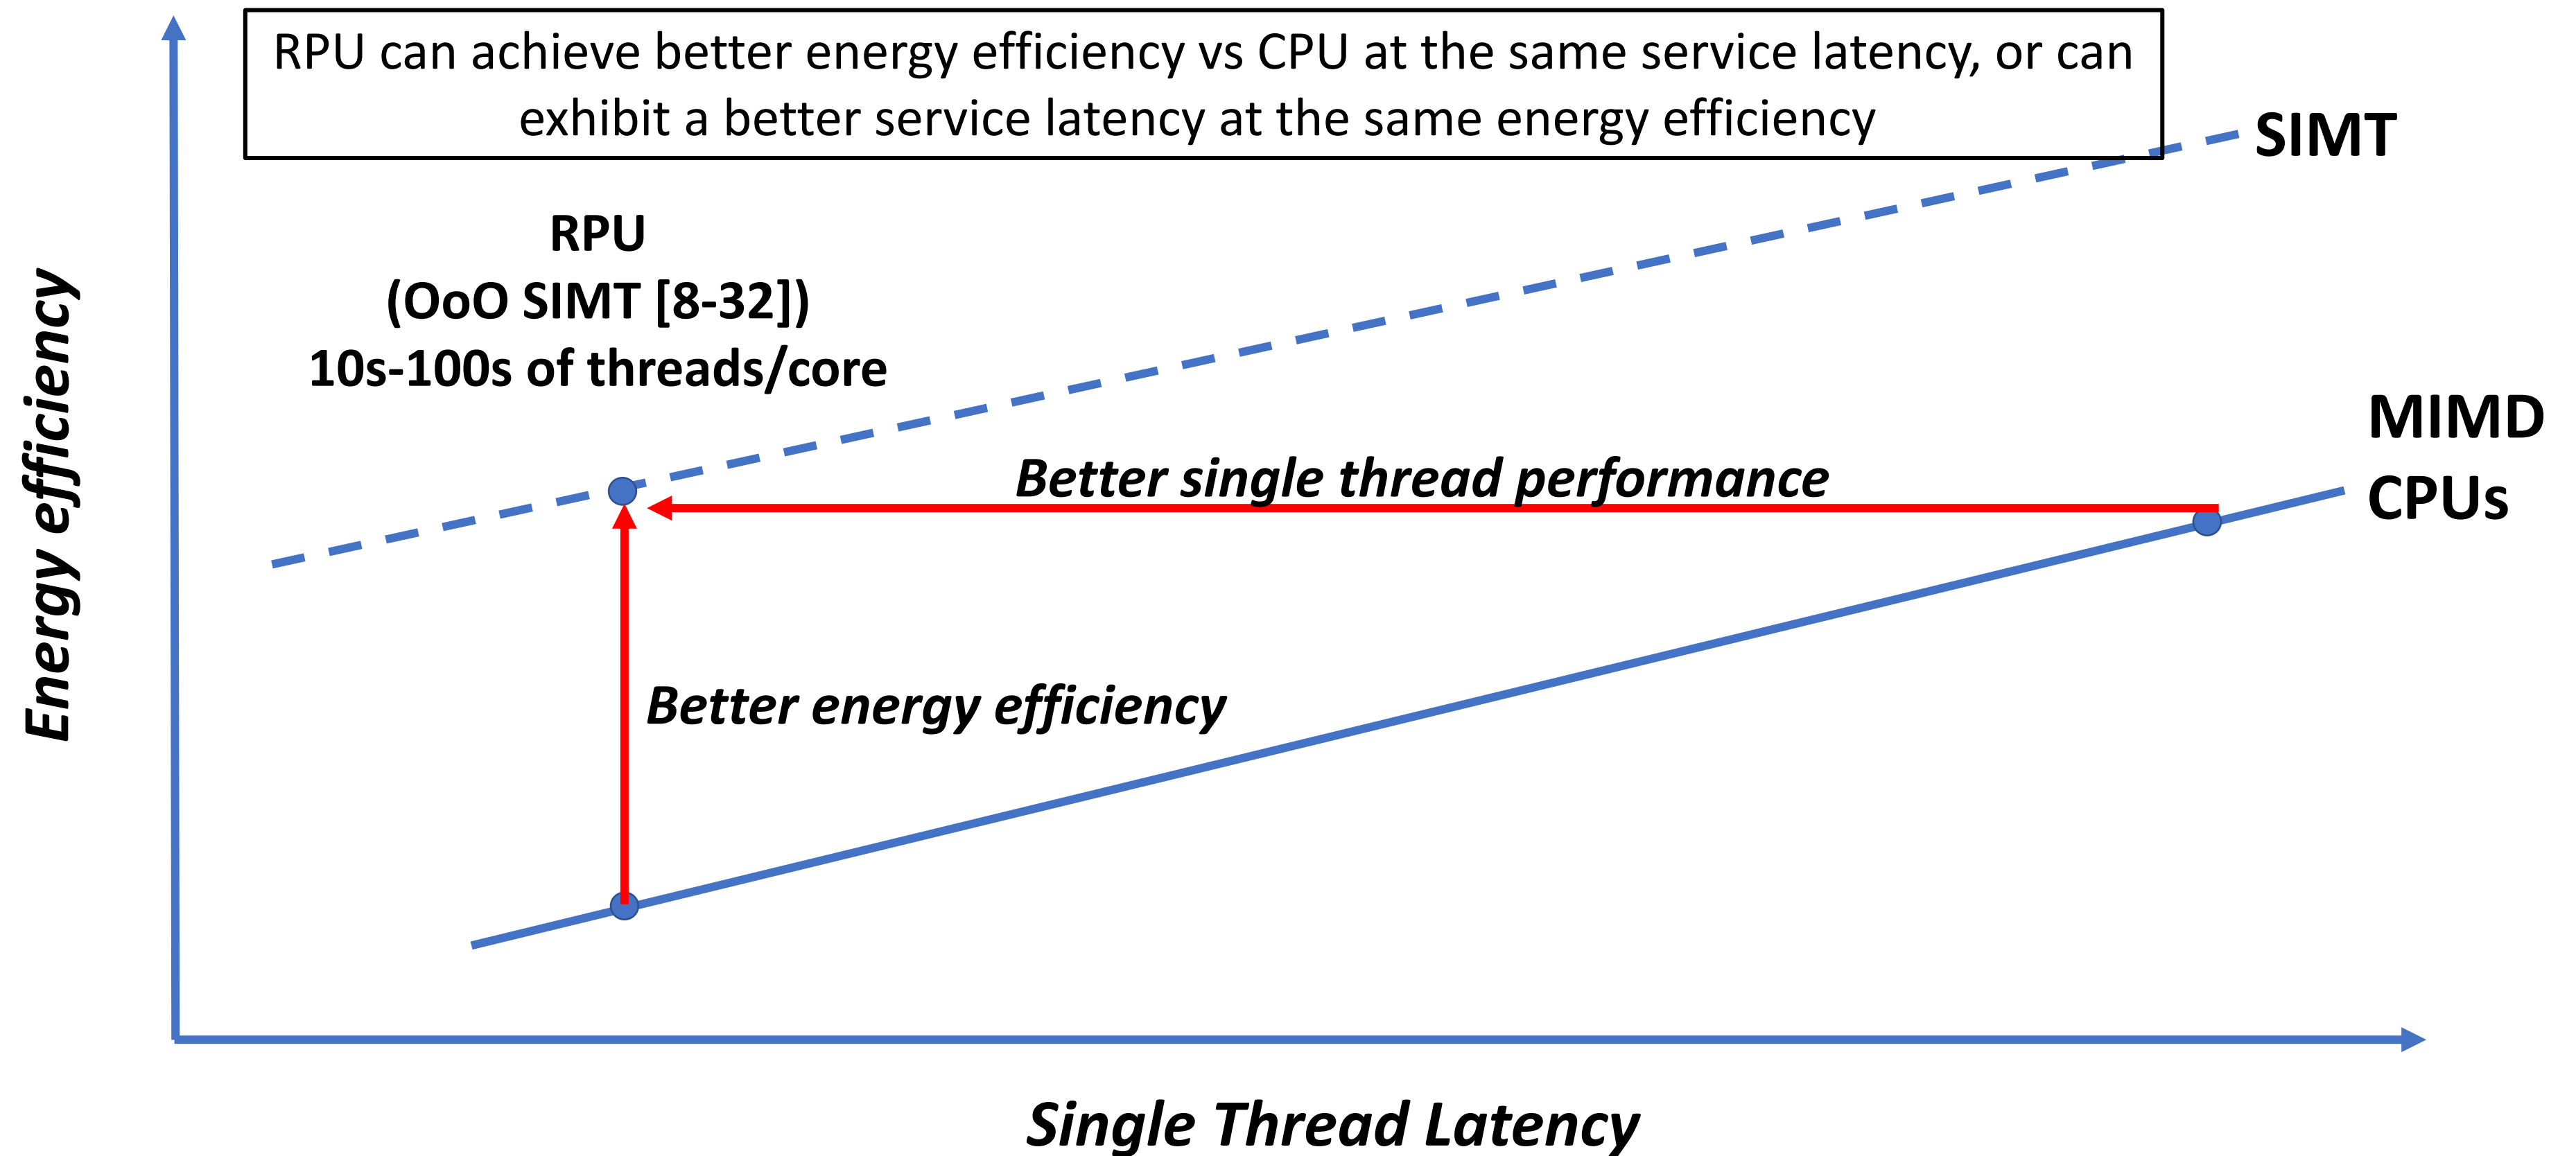
**Key Observation#4:** Modern data centers already rely on request batching heavily



# Latency & Energy-Efficiency Tradeoff



# Latency & Energy-Efficiency Tradeoff



# HW/SW Stack

<b>Webservice (C++, PHP, ...)</b>
<b>ARM/x86 compiler</b>
<b>HTTP server</b>
<b>Runtime/libs (pthread, cstdlib, ..)</b>
<b>OS (Process, VM, I/Os)</b>
<b>Multi Core CPU</b>

**CPU SW Stack**

<b>CUDA</b>
<b>CUDA compiler</b>
<b>Nvidia Triton HTTP server</b>
<b>CUDA runtime/libs (cudalib, tensorRT, ..)</b>
<b>OS (I/Os management)</b>
<b>CUDA driver (VM/thread management)</b>
<b>GPU Hardware</b>

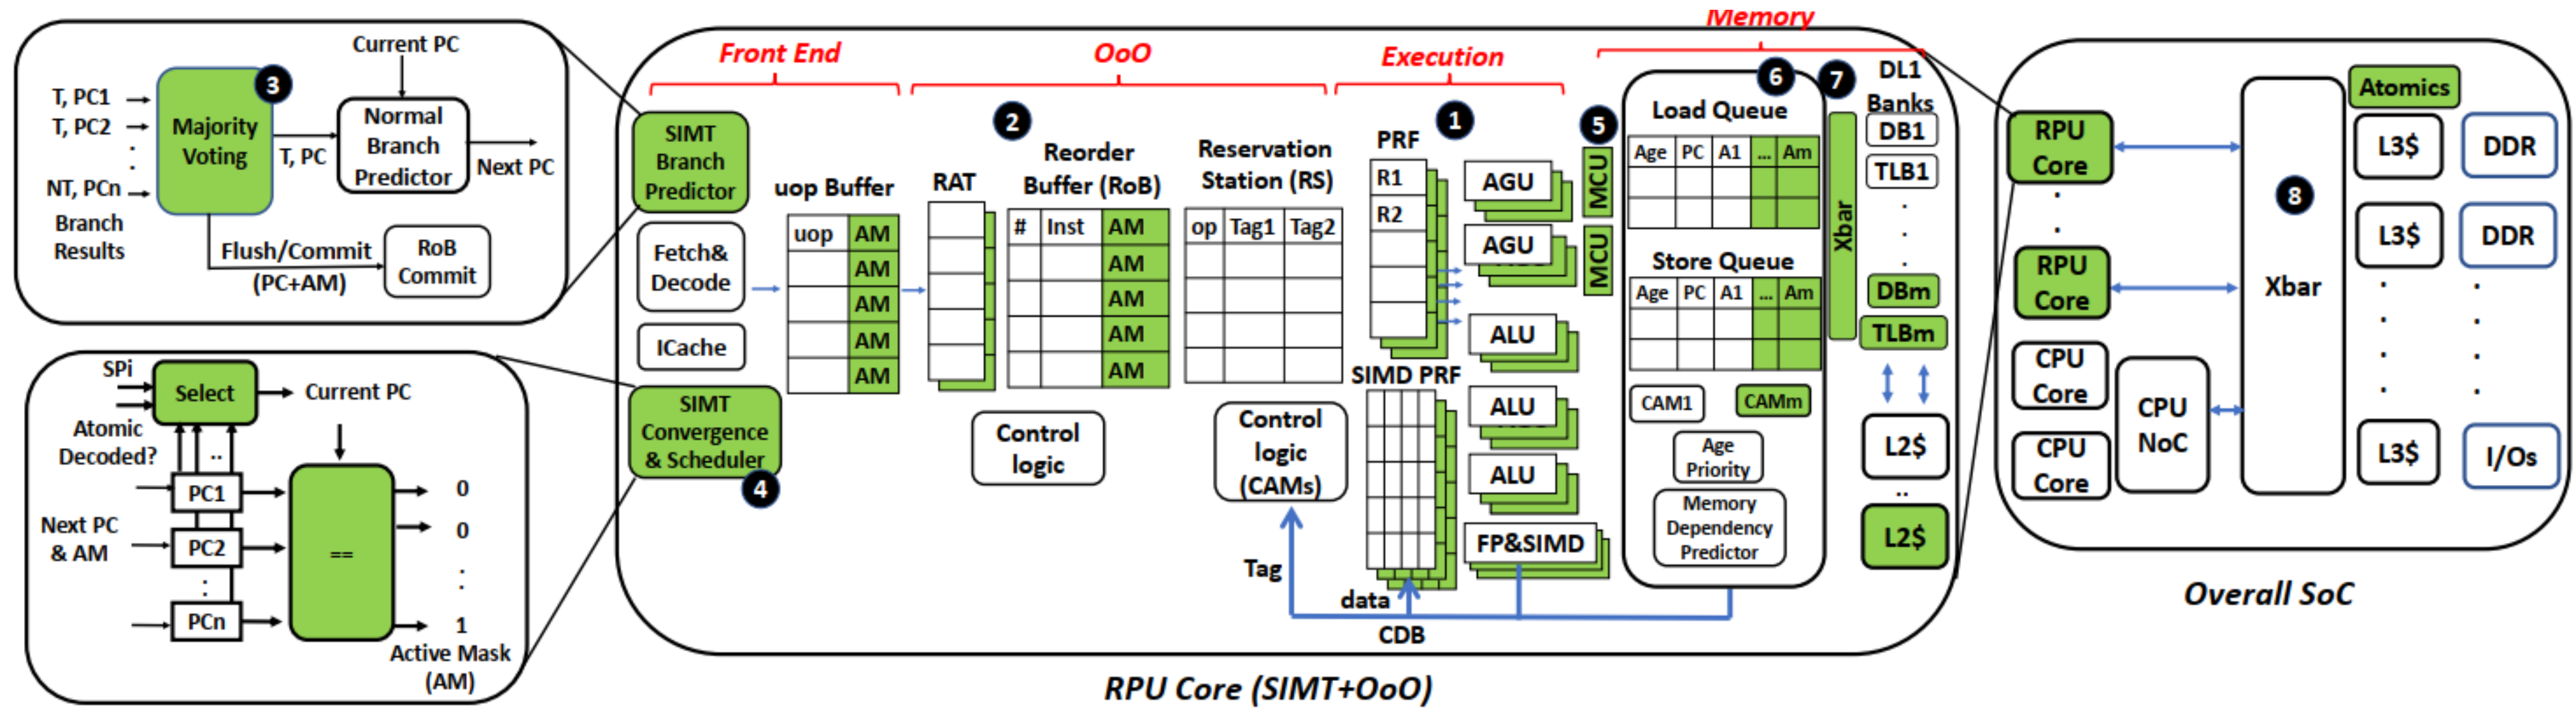
**GPU SW Stack**

<b>Webservice (C++, PHP, ...)</b>
<b>ARM/x86 compiler</b>
<b>Batch-aware HTTP server</b>
<b>Runtime/libs (pthread, cstdlib, ..)</b>
<b>OS (I/Os management)</b>
<b>RPU driver (VM/thread management)</b>
<b>RPU Hardware</b>

**RPU SW Stack**

- For RPU, we keep the SW programming interface as in the CPU
- Some VM&process management system calls are reimplemented in the RPU driver to be batch-aware

# RPU HW



# Energy Efficiency of CPU vs RPU (Analytical Model)

$$\frac{\text{CPU Energy}}{\text{RPU Energy}} = \frac{\text{Execution Energy} + \text{Memory system Energy} + \text{Front\_OoO Energy} + \text{Static Energy}}{\text{Execution Energy} + (1 - r) (\text{Memory system Energy}) + \frac{1}{n * \text{eff}} [\text{Front\_OoO Energy} + r * \text{Memory system Energy} + \text{Static Energy}] + \text{SIMT\_Overhead}}$$

*batch size (n) = 8-32*

*SIMT Efficiency=92%*

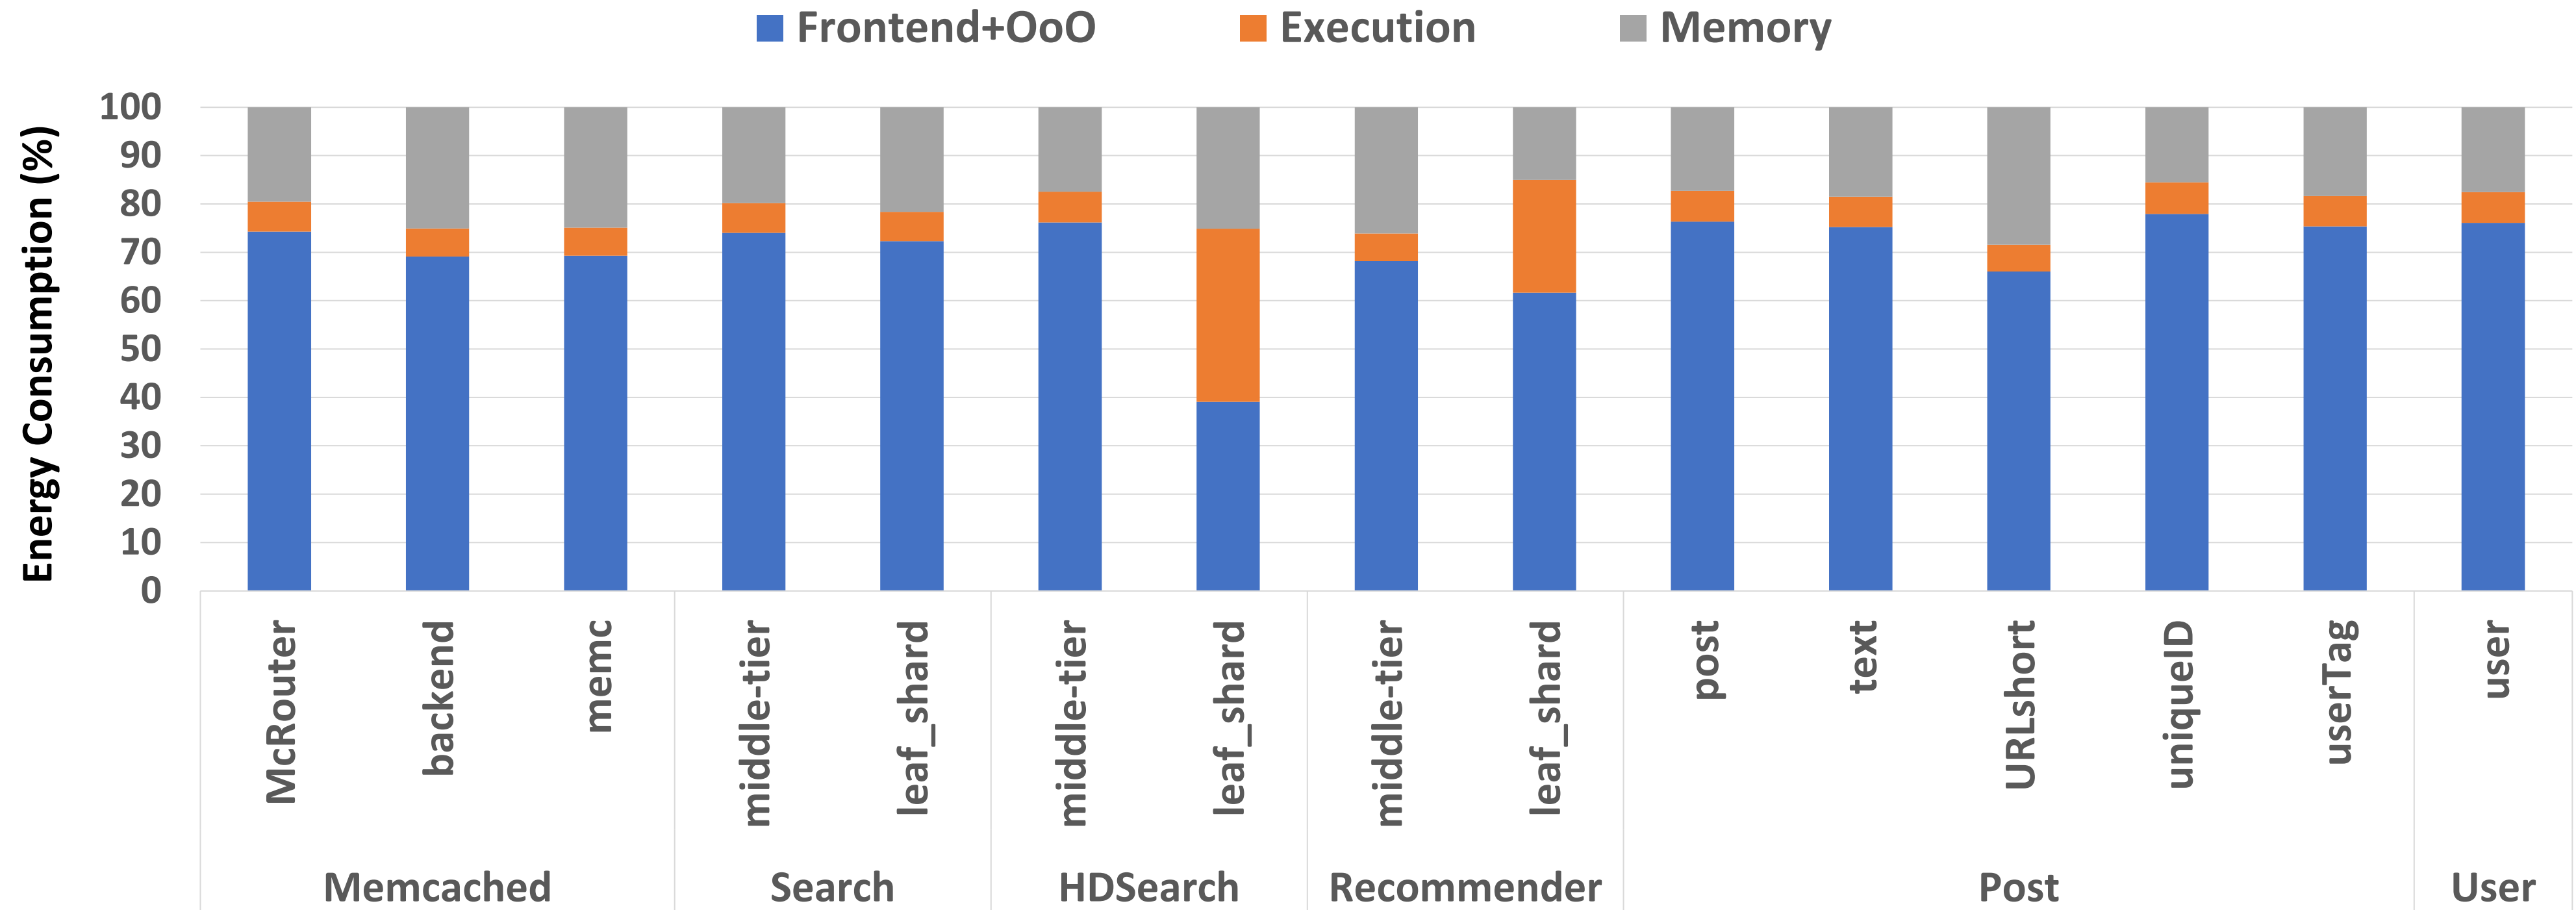
*data locality ratio =75%*

*Amortized factors = 50-90%*

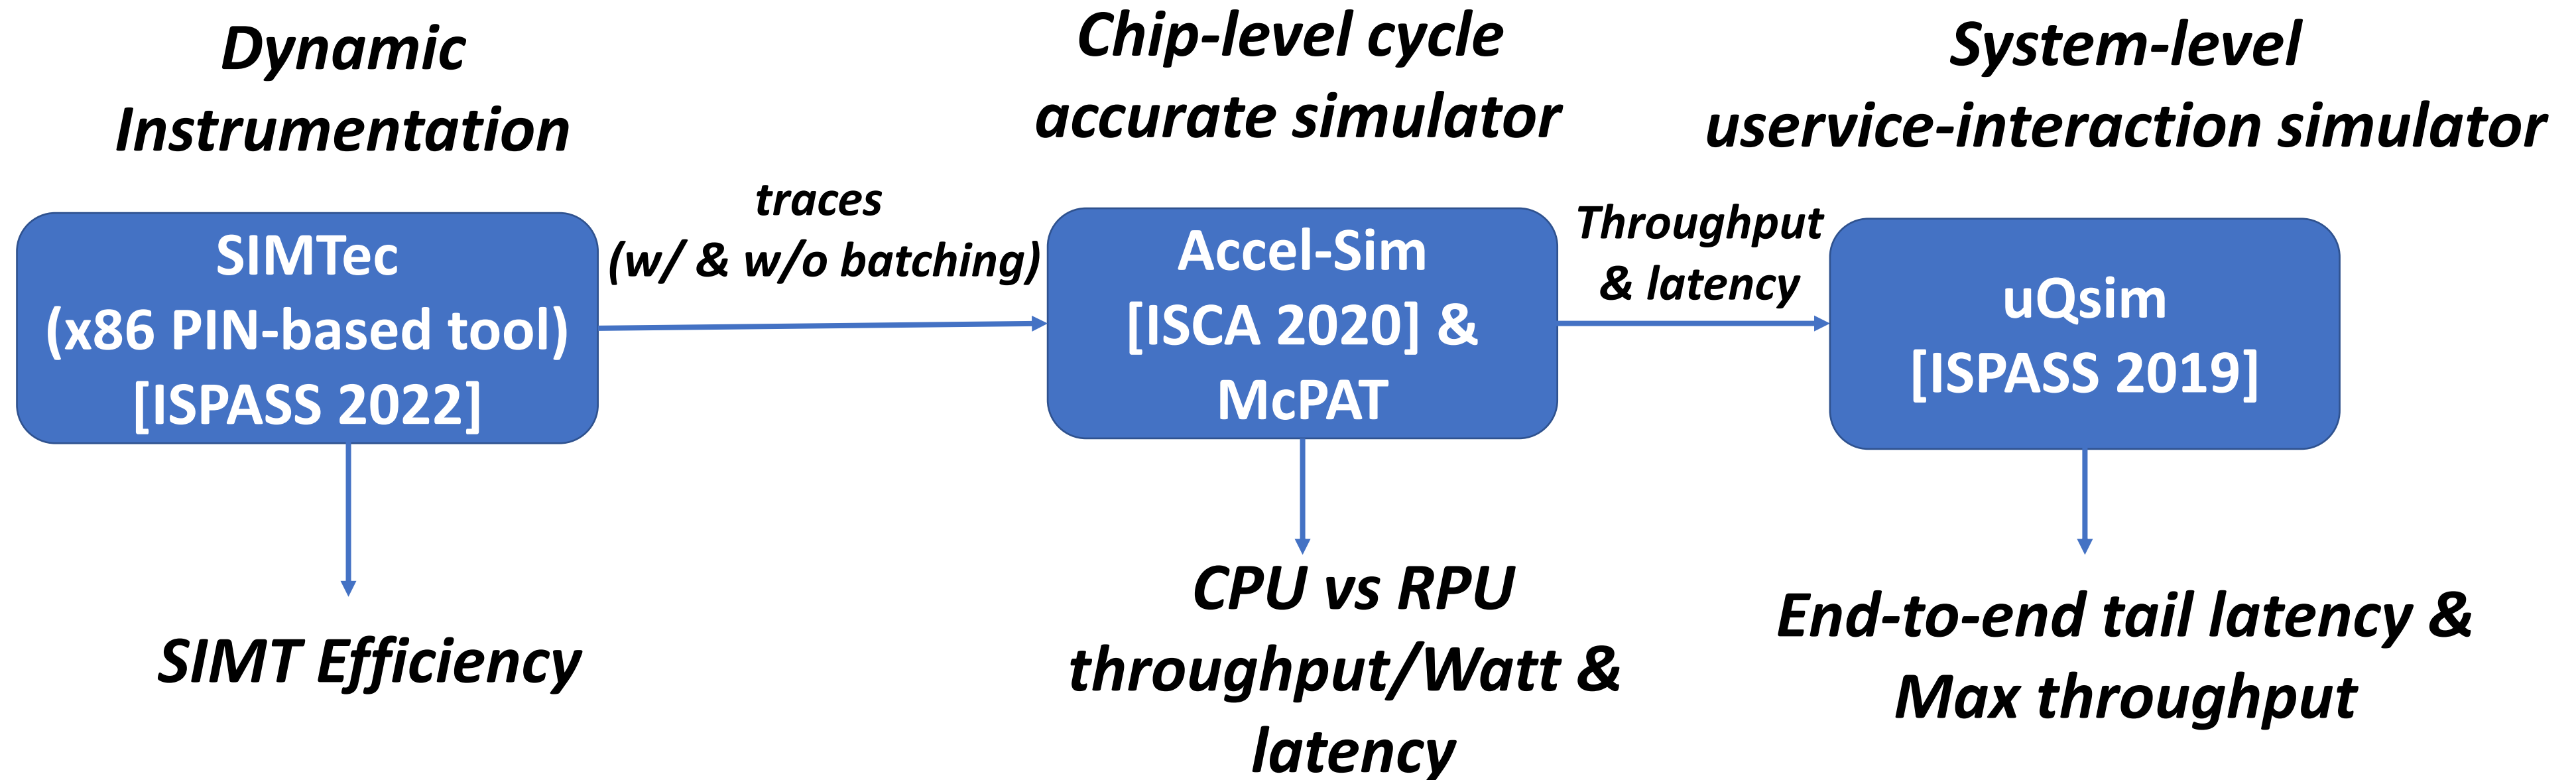
*Larger L1/L2 MCUs  
Active mask  
etc.*

→ an anticipated 2-10x energy efficiency gain can be achieved with RPU vs CPU

# CPU Dynamic Energy Breakdown



# Experimental Setup



## **Workloads: Social Network Microservices**

Microsuite [IISWC 2018], DeathStarBench [ASPLOS 2020] and In-house benchmarks

Libraries: c++ stdlib, Intel MKL, OpenSSL, FLANN, Pthread, zlib, protobuf, gRPC and MLPack, ...



# Batching Opportunity for Facebook Services

- To amortize batching overhead, you either need:
  - (1) High service latency, with low traffic so service latency will amortize batching **OR**
  - (2) High traffic, with low service latency so high traffic will amortize batching **OR**
  - (3) High traffic and high service latency (ideal case)
- Let's take a look at Facebook in-production services:

$\mu$ service	Throughput (QPS)	Req. latency	Insn./query	
Web	O (100)	O (ms)	O ( $10^6$ )	} <b>Low traffic but high latency</b>
Feed1	O (1000)	O (ms)	O ( $10^9$ )	
Feed2	O (10)	O (s)	O ( $10^9$ )	
Ads1	O (10)	O (ms)	O ( $10^9$ )	
Ads2	O (100)	O (ms)	O ( $10^9$ )	
Cache1	O (100K)	O ( $\mu$ s)	O ( $10^3$ )	} <b>Low latency but high traffic</b>
Cache2	O (100K)	O ( $\mu$ s)	O ( $10^3$ )	

Note: I was not able to calculate the exact batching overhead as the exact numbers are not shown and SLA (P99 latency) is not specified.

# Batching Opportunity for Google Services

- (1) From Google in-production ML inference services:
  - Batching is widely used for DL inference with size = 8-20 reqs based on traffic and latency

<i>Production</i>						<i>MLPerf 0.7</i>		
<i>DNN</i>	<i>ms</i>	<i>batch</i>	<i>DNN</i>	<i>ms</i>	<i>batch</i>	<i>DNN</i>	<i>ms</i>	<i>batch</i>
MLP0	7	200	RNN0	60	8	Resnet50	15	16
MLP1	20	168	RNN1	10	32	SSD	100	4
CNN0	10	8	BERT0	5	128	GNMT	250	16
CNN1	32	32	BERT1	10	64			

Table 5. Latency limit in ms and batch size picked for TPUv4i.

Quoted: “Clearly, datacenter applications limit latency, not batch size. Future DSAs should take advantage of larger batch sizes”

- (2) Further, Google search service has a high service latency (~10s ms) and high traffic (~100K QPS), so they are a good candidate for batching